

# Social Media Analysis

## Data Clan

Daniel Chen, Ghazal Erfani, Shiven Shashidhar

95851: Making Products Count: Data Science for Product Managers

# DSPM Project

<b>Introduction</b>	<b>3</b>
<b>Data Preparation</b>	<b>4</b>
<b>Exploratory Data Analysis</b>	<b>4</b>
<b>Modeling Approach</b>	<b>6</b>
<b>Results and Evaluation</b>	<b>7</b>
<b>Conclusion and Recommendations</b>	<b>11</b>
<b>Future Improvements</b>	<b>11</b>
<b>References</b>	<b>12</b>
<b>Additional Plots</b>	<b>12</b>

# Introduction

Today, almost all smartphones run on one of two operating systems; Google's Android or Apple's iOS. These two platforms accounted for more 99.7 percent of all new smartphones shipped in 2018. It's no surprise that these mobile operating systems dominate the market: they are both excellent. They have much in common with one another, as well as some key differences.

Given their market domination, it comes to no surprise that the battle of the two operating systems has raged for more than a decade now. Their quality, price, and value have been compared widely. This battle was most exemplified during the release of the iPhone 8/X and Samsung Galaxy S8 in 2017. These smartphones have been regarded by some as 'the most beautiful phones' ever made. They have also been consistently pitted against one another. Their designs, displays, special features, and values have been analyzed closely by many new sources and professional reviewers.

Ultimately, it is not the opinion of news sources or professional reviewers that determines the success of these smartphones. It's the opinion of the consumers. Thanks to the pervasiveness of social media, consumers' opinions about these phones are now widely available. One can analyze their social media posts, comments, and reviews from sources such as blogs, Facebook, Instagram, and Twitter to better understand the sentiments around these smartphones.

In this analysis, we are focused on answering two primary questions. Firstly, we would like to get a grasp of consumers' sentiments about the Galaxy S8 and iPhone 8/X prior and post their releases. Secondly, we would like to understand what attributes of the Galaxy S8 and iPhone 8/X are most important to consumers and how consumers feel about those attributes.

Answers to such questions can be of significant value to the creators of these devices: Apple and Samsung. They can use such sentiment analysis to gauge consumers' interests and questions, make better product and marketing decisions, and forecast sales. Additionally, such sentiment analysis can greatly benefit new consumers to the market. They can learn from the experiences and opinions of consumers similar to themselves and make more informed decisions about the devices they choose to purchase.

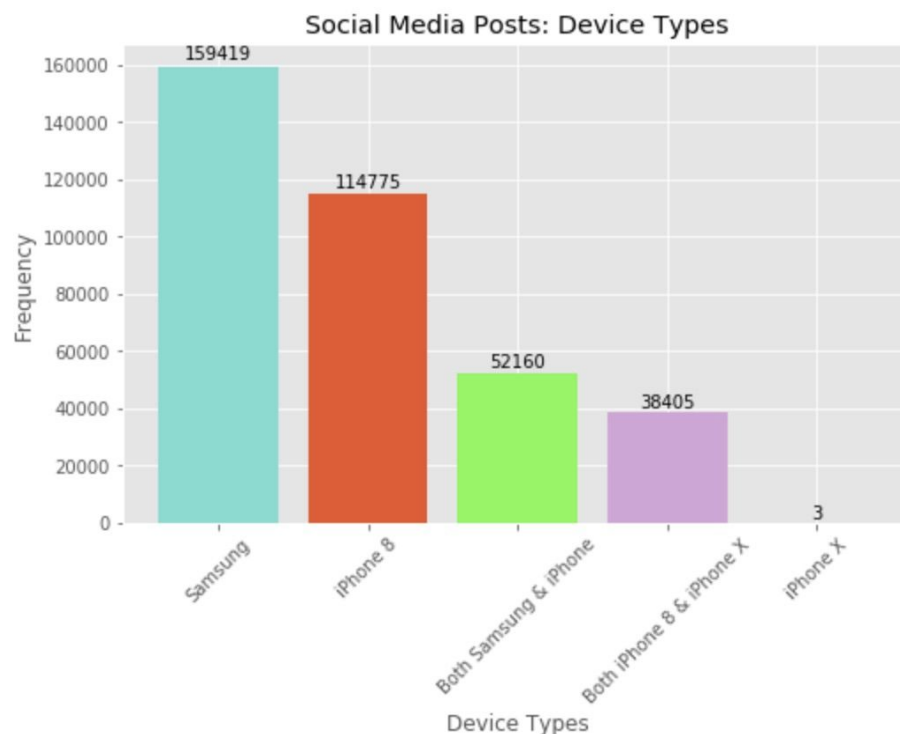
# Data Preparation

To prepare this data for modeling, we performed a series of data cleansing and extraction steps. Firstly, we removed all non-original posts, such as replies or comments, because these posts can be hard to contextualize since they are detached from their original post. Secondly, we eliminated any posts from authors with more than 100 unique followers to ensure we are only capturing end consumers, not news sources or professional reviewers. We then grouped the posts into ones that contained the device names, such as 'Galaxy S8.' We created the following five groups of device types: Galaxy S8, Galaxy S8 & iPhone, iPhone 8, iPhone X, iPhone 8 and iPhone X.

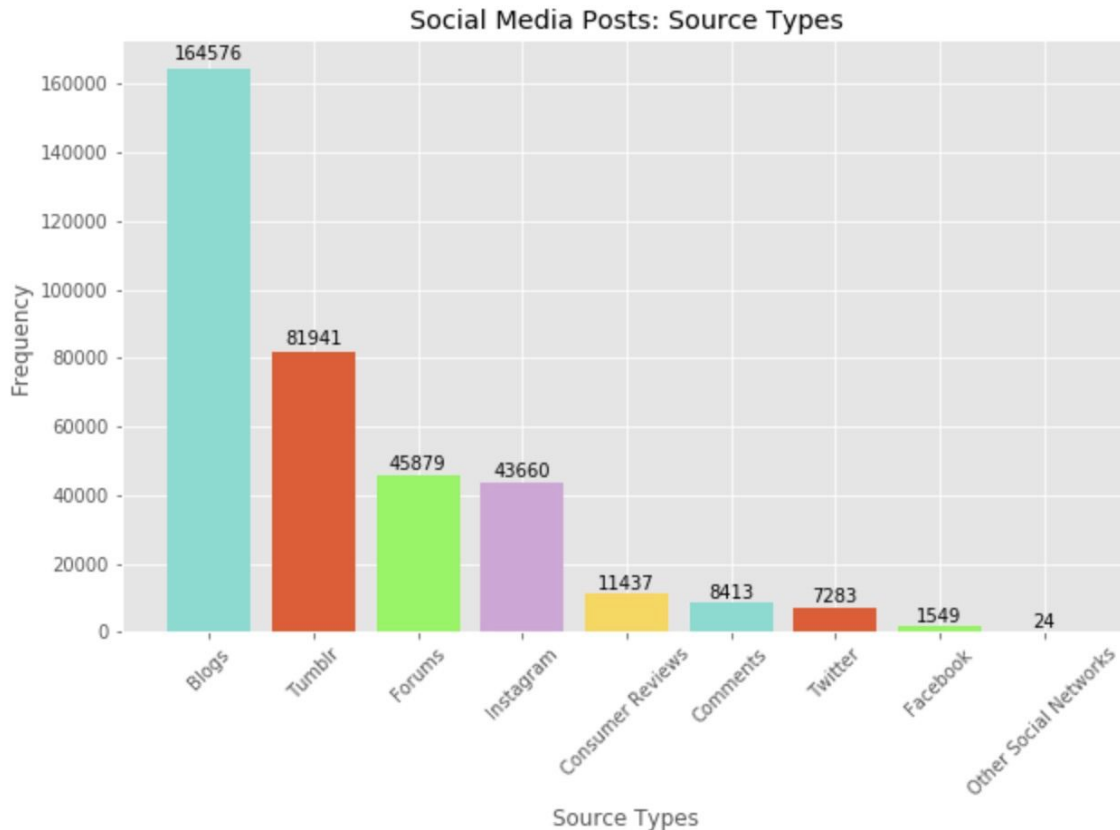
Next, we focused on preparing the posts for analysis. We removed punctuation and stopwords, converted to lowercase, tokenized and stemmed the posts, and did speech tagging. This allowed us to move forward with text analysis of the posts.

## Exploratory Data Analysis

To better understand the data, we performed exploratory data analysis. Firstly, we looked at the distribution of posts across device types.



It seemed that Samsung had the greatest number of posts while iPhone X had very few. Moving forward, we chose to not focus on the analysis of iPhone X only. We also looked at the spread of posts across different sources.



We learned that a large majority of the posts were from blogs. We were surprised to find that very few posts originated from social media sites such as Twitter and Facebook. We considered that a larger number of blog posts likely implied better text analysis, given that such text is often more structured than that in social media posts.

In addition, we performed topic modeling using Latent Dirichlet Allocation model to further explore the data. This was based on the following hypothesis: each user's post contains latent 'topics' (themes) and each topic is a distribution over different words/tokens. In our model, we considered each text post to be a probability distribution over 10 different 'topics' and each topic to be a distribution over 3324 words (in case of Apple iPhone) and 5893 words (in case of Samsung Galaxy). Here each 'topic' of a user review can be thought of as the main focus of the review. For example, the user might be commenting about the design of iPhone or the price of Samsung Galaxy. Then, the topic for the first case would be 'design' and for the second case

would be 'price.' Topic modeling enabled us to better understand the themes that are most pertinent to consumers.

Most significantly, we found that iPhone 8/X posts revealed four relevant topics: The first topic was primarily focused on the color, warranty and memory space. The second topic was focused on price/cost of the iPhone. The third topic mainly dealt with display, design, looks and screen quality. The fourth topic was based on wireless charging and power. Similarly, for Samsung Galaxy, we found four relevant topics: color, accessories, display and screen, and bixby and iris scanner. This provided us with some initial insights into consumers' primary concerns.

## Modeling Approach

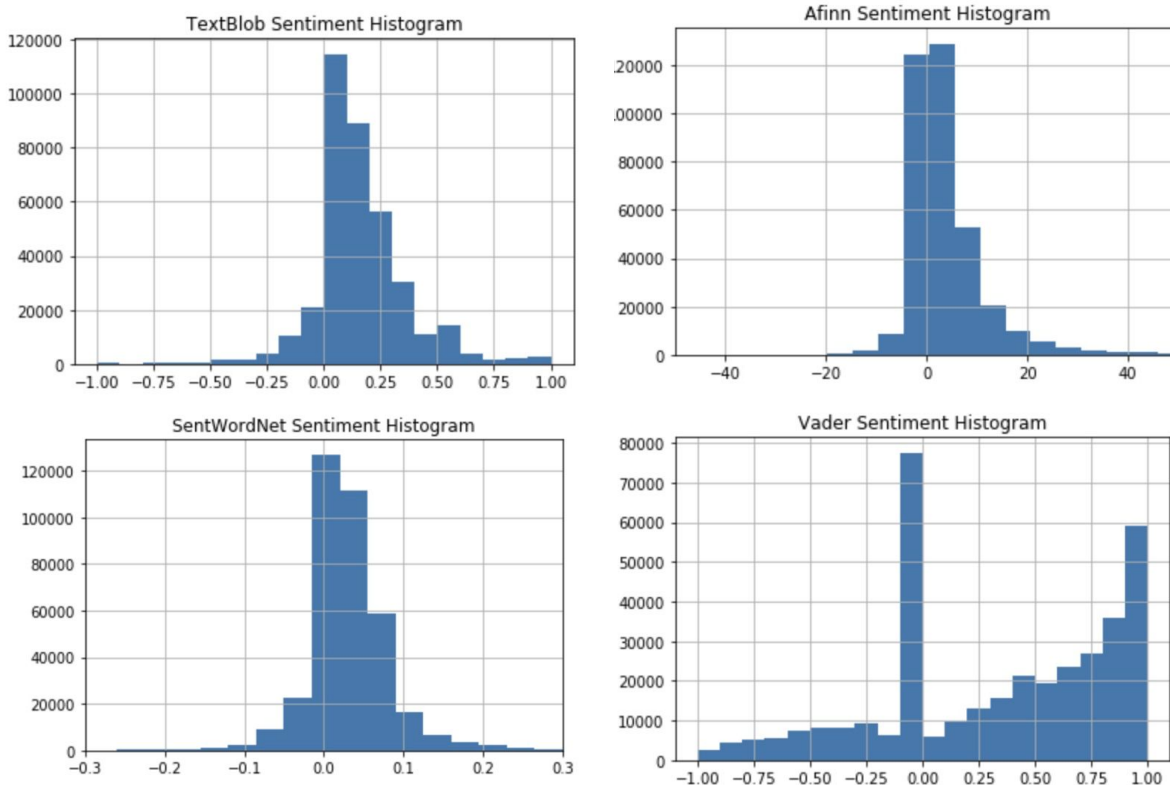
In aiming to better understand consumers' sentiments about the Galaxy S8 and iPhone 8/X prior and post their releases, we applied several sentiment analysis models. These robust, pre-trained models are tuned to analyze the sentiments of textual content and return a sentiment rating value that indicates negative, neutral, or positive. Given the quintillion bytes of data that are generated every day, sentiment analysis tools such as these are used ubiquitously to make sense of unstructured information and understand opinions about a given subject. We chose to use four models for our analysis: TextBlob, Afinn, SentiWordNet 3.0, and VADER (Valence Aware Dictionary and Sentiment Reasoner). Some of these models, such as TextBlob and SentiWordNet 3.0 are well-suited for any text analysis while others, such as Afinn and VADER are specifically attuned to sentiments expressed in social media, such as punctuation and emoticons. Given that we do not have labels to evaluate our models, we chose to apply all four models in our analysis to have a framework for comparison of sentiments across models.

In aiming to understand what attributes of the Galaxy S8 and iPhone 8/X are most important to consumers, we first used bigram analysis. As mentioned in the book 'Mining the Social Web', bigram analysis is a simple and powerful way for clustering commonly co-occurring words from social media and blog text data. A quick review of different mobile phone features revealed an interesting pattern: many product features occur as bigrams. For example, consider iPhone X product features as outlined on the 'Boost Mobile' website for iPhone X: TrueDepth Camera, Intuitive Gestures, OLED Display, Portrait Lighting to name a few. Hence, we primarily utilized bigram analysis to determine the most important product attributes for Apple iPhone and

Samsung Galaxy. We then fed these most important features to our sentiment analysis models to understand how consumers of each device feel about these features.

## Results and Evaluation

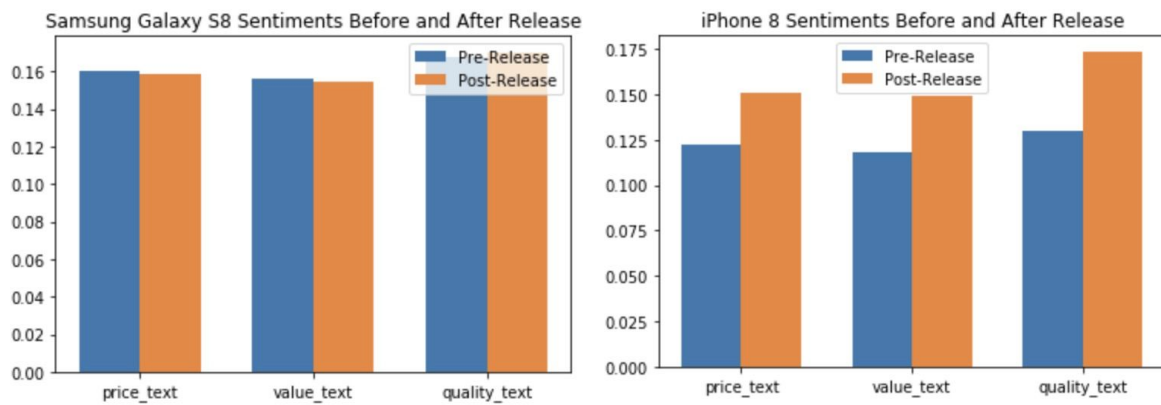
Firstly, we examined and compared the results of the four sentiment analysis models. As displayed by the graphs below, TextBlob, Afinn, and SentiWordNet models resulted in similar distributions of sentiments across all posts. The Vader model resulted in sentiments skewed towards the positive side.



In deciding which model to move forward with, we looked at several specific Twitter posts--due to their short length--and manually compared the sentiments of each model to find which model provided the most accurate one. With this manual analysis, we found that TextBlob seemed to provide the most accurate results and chose to move forward with this model.

Next, we used the results of the sentiment analysis to understand consumers' sentiments around quality, price, and value of the devices. To do this, we identified posts related to quality,

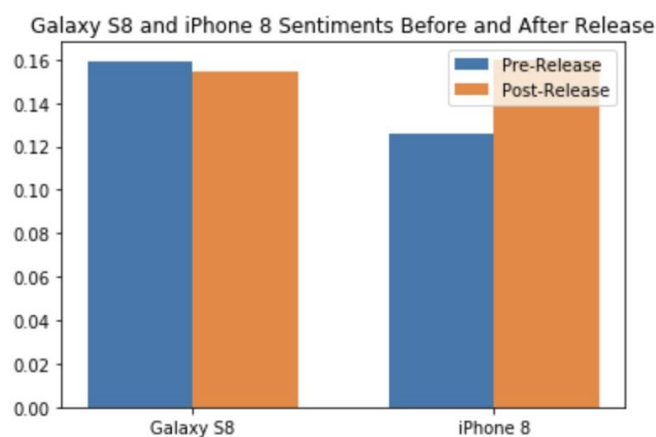
price, and value based on the words contained in each post and observed the sentiments of each post prior to and post release of each device. For Galaxy S8, we found that sentiments around quality, price, and value did not change much before and after its release. However, for iPhone 8, we found that sentiments around these three topics improved dramatically, most significantly around



quality.

Overall, sentiments around iPhone 8 prior to its release were a bit lower than sentiments around Galaxy S8 prior to its release. However, sentiments for iPhone 8 became significantly more positive after its release whereas sentiments for Galaxy S8 became slightly more negative.

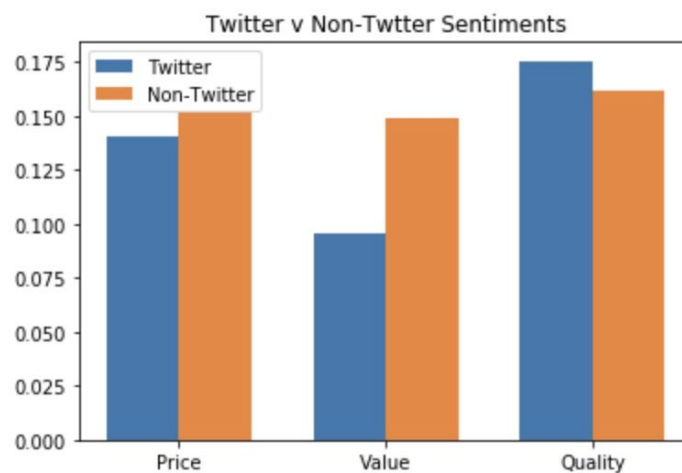
Given that important data source we chose to explore quality, price, and posts compared to



Twitter is an important data source for this analysis, sentiments around value of Twitter non-Twitter posts.



We found that sentiments on Twitter were more positive around value and price and more negative around quality compared to non-Twitter posts. This may be because Twitter posts are primarily opinions of direct consumers who may be more critical about price and the value they get out of the phone day-to-day. In contrast, non-Twitter posts, such as blogs, may be more critical of the long-term quality of phone features, not necessarily the day-to-day value to the consumer.



Next, to get a good picture of what features of the products are being discussed by the users the most, we looked at the 30 most frequently occurring bigrams in the reviews/post based on the social media analysis of Apple and Galaxy S8 posts. These features can be considered as the most important product attributes due to their frequent citations in posts. The following features were the most important attributes for Apple iPhone users in descending order:


1. Wireless charging
2. Home button (Apple removed home button from iPhone X, so this generated a good amount of discussion among users)
3. Fast Charge
4. Portrait lighting (New feature introduced in iPhone)
5. Bionic Chip (Apple introduced this new feature iPhone 8 onwards)
6. Facial recognition (Apple introduced this new feature iPhone 8 onwards)
7. Charge Pad (Apple introduced this new feature iPhone 8 onwards)
8. Dual Camera

## 9. Oled Display (New OLED screen for iPhone X)

The following features were the most important attributes for Samsung Galaxy users in descending order:

1. Memory (as indicated by 'gb ram')
2. Infinity display
3. Iris Scanner
4. Bixby button
5. Fingerprint sensor/scanner
6. Bixby voice
7. Facial recognition
8. Battery life
9. Wireless charger
10. Red tint (an issue faced by some users with Samsung products)

To understand users' sentiments around these most important features, we applied the sentiment analysis to posts that contained content around these features.



We found that sentiments are generally more positive for iPhone 8 than those for Galaxy S8. In particular, there were more positive sentiments around iPhone 8 accessories, but more positive sentiments around Galaxy S8 battery life.

## Conclusion and Recommendations

From a product manager's perspective, we can look at the data and get a strong idea about what features are most important to the users of our device. Once we have determined a feature to look at it is possible to analyze the distribution of how people feel about that specific feature. Using this knowledge, product managers can better allocate resources to develop or fix aspects of the product that is subpar. Also, knowing the strengths of a product from a user perspective is a valuable insight for a marketing team to help them determine what features of the product to promote to potential users.

This analysis shows that Samsung's battery life is a great part of the device and is driving much of what users are talking about. People also seem to really enjoy the dual camera. However, they are lacking in the repair and maintenance aspect of their business. This could be a good area of focus for a product manager.

Apple has been doing a great job when it comes to accessories and complementary aspects of their product. From the bi-gram analysis, we can see that apple watch, apple tv, and the apple store are all things that are very important to their users. Implementing more features that allow Apple devices to connect seamlessly to each other should be a priority. On the other hand, there are a few areas where Apple can improve so they can be more competitive in the market. Facial recognition and wireless charging are two features that rank highly among what consumers want and where iPhone 8 is falling short of their competitor.

## Future Improvements

The methodology we used for analyzing tweets was identical to both blogs and longer articles. We assigned one sentiment score over for each review/tweet no matter how long the article was. An improvement would be to segment the longer articles so we can try and isolate parts of an article that are relevant to a specific feature and do analysis just on the relevant portion.

We also would like to implement our own sentiment analyzer by creating an LSTM RNN. A custom RNN would be useful to more accurately analyze review and text that are in this specific domain. While the out of the box sentiment analyzers performed fairly well, they were not designed specifically with our corpus in mind and has room for improvement. We would model this approach after SES, a self-supervised syntax based method of classification, that has been developed from Peking University.


## References

1. <https://www.digitaltrends.com/mobile/android-vs-ios/>
2. <https://www.gamingscan.com/ios-vs-android/>
3. <https://www.tomsguide.com/us/iphone-x-vs-galaxy-s8,review-4864.html>
4. <https://monkeylearn.com/sentiment-analysis/>
5. <https://www.webpages.uidaho.edu/~stevel/504/mining-the-social-web-2nd-edition.pdf>
6. <https://textblob.readthedocs.io/en/dev/>
7. <https://medium.com/analytics-vidhya/simplifying-social-media-sentiment-analysis-using-vader-in-python-f9e6ec6fc52f>
8. [http://corpustext.com/reference/sentiment\\_afinn.html](http://corpustext.com/reference/sentiment_afinn.html)
9. <http://nmis.isti.cnr.it/sebastiani/Publications/LREC10.pdf>
10. <http://aclweb.org/anthology/Y09-2018>

## Additional Plots





## iPhone 8/X Word Cloud



## Galaxy S8 Word Cloud







# 95-851: Social Media Analysis

## Data Clan

### Part 1: Sentiment Analysis

#### 1. Load Data

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
plt.style.use('ggplot')
%matplotlib inline
```

```
In [2]: Both_iPhones_data = pd.read_csv('Both_iPhones_data.csv')
both_samsung_and_iphone_data = pd.read_csv('both_samsung_and_iphone_data.csv')
iphone8_data = pd.read_csv('iphone8_data.csv')
iphoneX_data = pd.read_csv('iphoneX_data.csv')
samsung_data = pd.read_csv('samsung_data.csv')
```

#### 2. Exploratory Data Analysis

```
In [3]: # add new column with device type
Both_iPhones_data['device_type'] = 'Both iPhone 8 & iPhone X'
both_samsung_and_iphone_data['device_type'] = 'Both Samsung & iPhone'
iphone8_data['device_type'] = 'iPhone 8'
iphoneX_data['device_type'] = 'iPhone X'
samsung_data['device_type'] = 'Samsung'
```

```
In [4]: # concatenate all data into one df
all_data = pd.concat([Both_iPhones_data, both_samsung_and_iphone_data, iphone8_data, iphoneX_data, samsung_data], sort=True)
```

```
In [5]: def autolabel(rects):
    # attach some text labels
    for rect in rects:
        height = rect.get_height()
        plt.text(rect.get_x() + rect.get_width()/2., 1.01*height,
                 '%d' % int(height),
                 ha='center', va='bottom')
```



```

In [6]: from collections import Counter
        from operator import itemgetter
        c=['turquoise', 'orangered', 'lime', 'plum', 'gold']

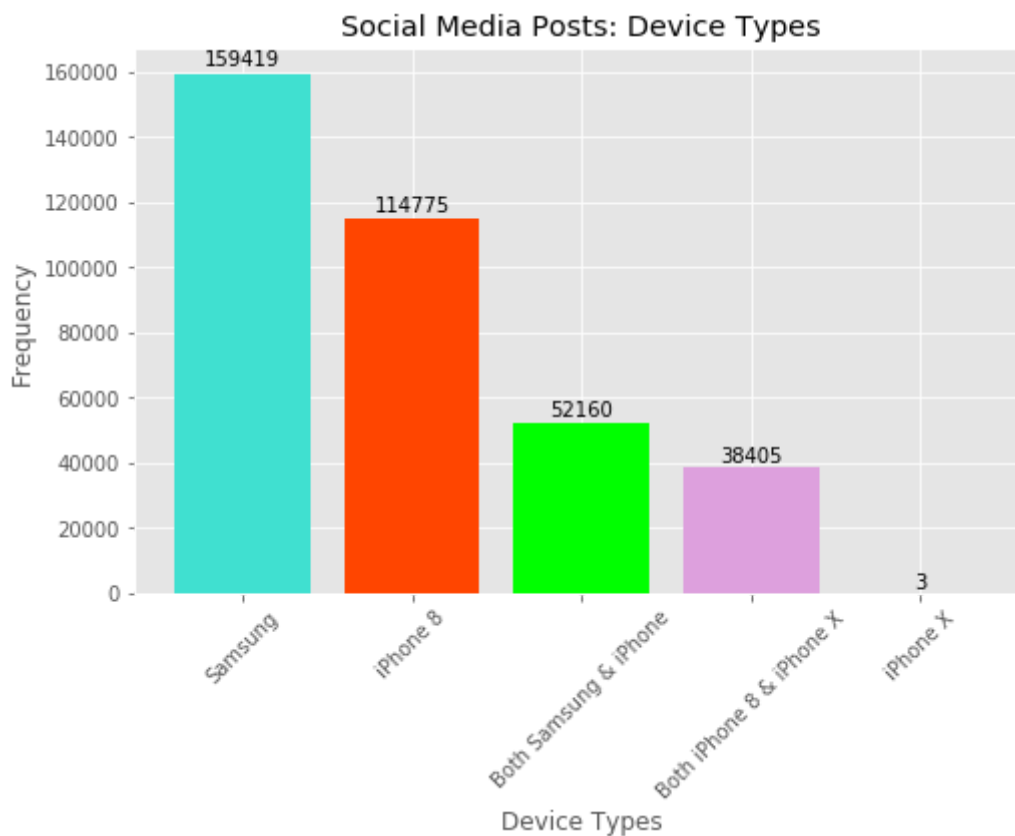
        counts_per_device_type = Counter()
        for device_type in all_data.device_type:
            counts_per_device_type[device_type] +=1

        counts_per_device_type_sorted = sorted(counts_per_device_type.items(),
                                                reverse=True,
                                                key=itemgetter(1))

        device_types = [device_type for device_type, count in counts_per_device_type_sorted]
        counts = [count for device_type, count in counts_per_device_type_sorted]

        plt.figure(figsize=(8, 5))
        bar1 = plt.bar(range(len(device_types)), counts, color = c)
        plt.xlabel('Device Types')
        plt.xticks(range(len(device_types)), device_types, rotation=45)
        plt.ylabel('Frequency')
        plt.title("Social Media Posts: Device Types")
        autolabel(bar1)
        plt.show()

```



```

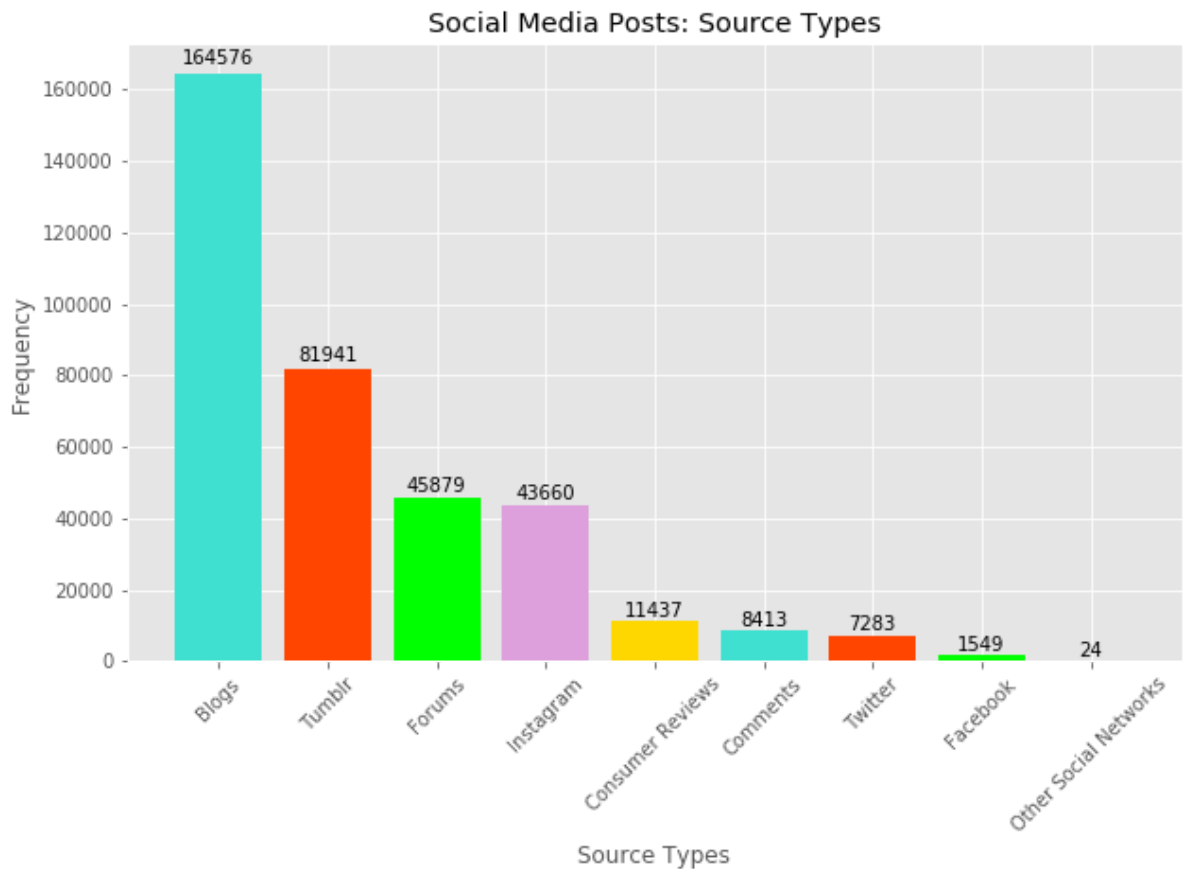
In [7]: from collections import Counter
from operator import itemgetter
c=['turquoise', 'orangered', 'lime', 'plum', 'gold']
counts_per_source_type = Counter()
for source_type in all_data["Source Type"]:
    counts_per_source_type[source_type] +=1

counts_per_source_type_sorted = sorted(counts_per_source_type.items(),
                                       reverse=True,
                                       key=itemgetter(1))

source_types = [source_type for source_type, count in counts_per_source_type_sorted]
counts = [count for source_type, count in counts_per_source_type_sorted]

plt.figure(figsize=(10, 6))
bar2 = plt.bar(range(len(source_types)), counts, color = c)
plt.xlabel('Source Types')
plt.xticks(range(len(source_types)), source_types, rotation=45)
plt.ylabel('Frequency')
plt.title("Social Media Posts: Source Types")
autolabel(bar2)
plt.show()

```



### 3. Data Preprocessing

```
In [9]: %%time
# remove punctuation
import string
all_data.LowerText = all_data.LowerText.apply(lambda x: x.translate(string.punctuation))
```

Wall time: 16.4 s

```
In [10]: %%time
# remove numbers
all_data.LowerText = all_data.LowerText.apply(lambda x: x.translate(string.digits))
```

Wall time: 15 s

```
In [11]: %%time
#tokenize
all_data["Tokenized_Text"] = all_data.LowerText.apply(lambda x: x.split(' '))
```

Wall time: 8.55 s

```
In [12]: %%time
# stem
import nltk
from nltk.stem.porter import PorterStemmer
all_data["Stemmed_text"] = all_data.Tokenized_Text.apply(lambda x: [PorterStemmer().stem(y) for y in x])
```

Wall time: 13min

```
In [15]: %%time
# remove stopwords
from nltk.corpus import stopwords
# import nltk
# nltk.download('stopwords')
sw = stopwords.words('english')
all_data.Stemmed_text = all_data.Stemmed_text.apply(lambda x: [item for item in x if item not in sw])
```

```
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\dpc50\AppData\Roaming\nltk_data...
[nltk_data] Unzipping corpora\stopwords.zip.
```

Wall time: 1min 12s

## 4. Create Sentiment Analysis Models

```
In [16]: %%time
from textblob import TextBlob
all_data['TextBlob_sentiment'] = np.array([TextBlob(text).sentiment.polarity
for text in all_data['Sound Bite Text']])
```

Wall time: 6min 30s

```
In [17]: %%time
from afinn import Afinn
afn = Afinn(emoticons=True)
all_data["Afn_Sent"] = all_data["Sound Bite Text"].apply(lambda x: afn.score(x))
```

Wall time: 21min 18s

```

In [50]: import spacy
nlp = spacy.load('en')
from nltk.corpus import sentiwordnet as swn

def analyze_sentiment_sentiwordnet_lexicon(review,
                                           verbose=False):

    # tokenize and POS tag text tokens
    tagged_text = [(token.text, token.tag_) for token in nlp(review)]
    pos_score = neg_score = token_count = obj_score = 0
    # get wordnet synsets based on POS tags
    # get sentiment scores if synsets are found
    for word, tag in tagged_text:
        ss_set = None
        if 'NN' in tag and list(swn.senti_synsets(word, 'n')):
            ss_set = list(swn.senti_synsets(word, 'n'))[0]
        elif 'VB' in tag and list(swn.senti_synsets(word, 'v')):
            ss_set = list(swn.senti_synsets(word, 'v'))[0]
        elif 'JJ' in tag and list(swn.senti_synsets(word, 'a')):
            ss_set = list(swn.senti_synsets(word, 'a'))[0]
        elif 'RB' in tag and list(swn.senti_synsets(word, 'r')):
            ss_set = list(swn.senti_synsets(word, 'r'))[0]
        # if senti-synset is found
        if ss_set:
            # add scores for all found synsets
            pos_score += ss_set.pos_score()
            neg_score += ss_set.neg_score()
            obj_score += ss_set.obj_score()
            token_count += 1

    # aggregate final scores
    final_score = pos_score - neg_score
    norm_final_score = 0
    if(token_count != 0):
        norm_final_score = round(float(final_score) / token_count, 2)

    final_sentiment = 'positive' if norm_final_score >= 0 else 'negative'

    if verbose:
        norm_obj_score = round(float(obj_score) / token_count, 2)
        norm_pos_score = round(float(pos_score) / token_count, 2)
        norm_neg_score = round(float(neg_score) / token_count, 2)
        # to display results in a nice table
        sentiment_frame = pd.DataFrame([[final_sentiment, norm_obj_score,
                                         norm_pos_score, norm_neg_score, norm_final_score],
                                         [0, 1, 2, 3, 4]])
        columns=pd.MultiIndex(levels=[['SENTIMENT STATS:'],
                                       ['Predicted Sentiment', 'Objectivity', 'Positive', 'Negative', 'Overall']],
                               labels=[[0, 0, 0, 0], [0, 1, 2, 3, 4]])
        print(sentiment_frame)

```

```
return norm_final_score
```

```
In [51]: %%time
all_data["Sentwordnet_sent"] = all_data["Sound Bite Text"].apply(lambda
x: analyze_sentiment_sentiwordnet_lexicon(x))
```


Wall time: 4h 15min 40s

```
In [21]: %%time
# vader model
from vaderSentiment.vaderSentiment import SentimentIntensityAnalyzer
sia = SentimentIntensityAnalyzer()
all_data['Vader_Sentiment'] = all_data['Sound Bite Text'].apply(lambda
x: float(sia.polarity_scores(x)['compound']))
```

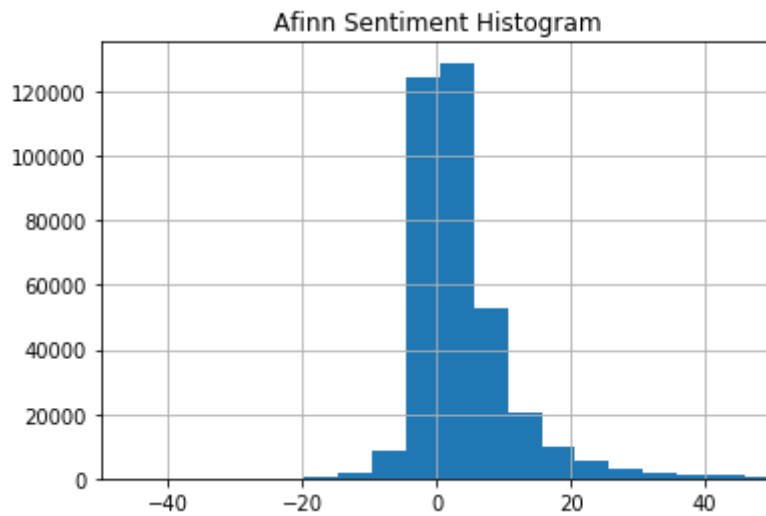
Wall time: 15min 3s

```
In [52]: all_data.to_csv("all_data_sent.csv", index = False)
```


```
In [139]: all_data.TextBlob_sentiment.hist(bins = 20)
plt.title("TextBlob Sentiment Histogram")
plt.show()
```



```
In [140]: all_data.Afn_Sent.hist(bins = 100)
plt.title("Afinn Sentiment Histogram")
plt.xlim(-50,50)
plt.show()
```




```
In [141]: all_data.Sentwordnet_sent.hist(bins = 50)
plt.title("SentWordNet Sentiment Histogram")
plt.xlim(-.3,.3)
plt.show()
```



```
In [142]: all_data.Vader_Sentiment.hist(bins = 20)
plt.title("Vader Sentiment Histogram")

plt.show()
```



```
In [117]: twitter_data = all_data[all_data["Source Type"] == "Twitter"]
twitter_data.shape
```

```
Out[117]: (7283, 22)
```

```
In [121]: twitter_df = pd.DataFrame(twitter_data)
twitter_df.columns = all_data.columns
twitter_df["sent_diff"] = abs(twitter_df.TextBlob_sentiment - twitter_df
.Sentwordnet_sent)
```

```
In [124]: sorted_twitter = twitter_df.sort_values(by = "sent_diff",ascending = False)
```



```
In [129]: for i in range(5):
            text = sorted_twitter.iloc[i,5]
            print(text)

            for y in range(14,18):
                print(sorted_twitter.columns[y],sorted_twitter.iloc[i,y])
            print()
```

I added a video to a @YouTube playlist [youtu.be/0MxC9sVp6mg?](https://youtu.be/0MxC9sVp6mg?) Apple Iph one 8 Ringtone Awesome (Never Ever)

TextBlob\_sentiment 1.0  
Afn\_Sent 4.0  
Vader\_Sentiment 0.6249  
Sentwordnet\_sent -0.1

Report: Galaxy S8 Prototypes Impressed At MWC 2017 [dlvr.it/NYHFyP](https://dlvr.it/NYHFyP) @slide me [pic.twitter.com/OW9WIvw2uX](https://pic.twitter.com/OW9WIvw2uX)

TextBlob\_sentiment 1.0  
Afn\_Sent 3.0  
Vader\_Sentiment 0.4767  
Sentwordnet\_sent -0.08

Apple Leak Reveals iPhone 8 Nasty Surprises [newssummedup.com/a/wpslfyfb.me/3gKqVChkT](https://newssummedup.com/a/wpslfyfb.me/3gKqVChkT)

TextBlob\_sentiment -1.0  
Afn\_Sent -4.0  
Vader\_Sentiment -0.6249  
Sentwordnet\_sent 0.03

Just received my galaxy S8 plus and one of my earphones don't work Sprint won't do anything. Service is terrible. @sprint #sprintservice

TextBlob\_sentiment -1.0  
Afn\_Sent 0.0  
Vader\_Sentiment -0.4767  
Sentwordnet\_sent 0.01

What's the Best #Samsung Galaxy S8 Case? 5 Affordable Ways to Protect Your Investment [muo.co/2s3YDH3](https://muo.co/2s3YDH3) [pic.twitter.com/8y49tQEIPJ](https://pic.twitter.com/8y49tQEIPJ)

TextBlob\_sentiment 1.0  
Afn\_Sent 6.0  
Vader\_Sentiment 0.7783  
Sentwordnet\_sent 0.0

```
In [118]: random_data = np.random.permutation/twitter_data)[:10]
            random_tweets_df = pd.DataFrame(random_data)
            random_tweets_df.columns = all_data.columns
```

```
In [119]:
```

In [120]:

Out[120]:

	LowerText	Media Type	Post ID	Post Type	Published Date (GMT-04:00) New York	Sound Bite Tr
3	i liked a @youtube video youtu.be/n2wjlnnrhi?...	Link	911756621767958530	Original	Sep 23, 2017 8:58:12 PM	I liked a @YouTube vid youtu.be/n2wJLnNJRHI'
7	i liked a @youtube video youtu.be/xkjblK52d1w?...	Link	854526508865646592	Original	Apr 18, 2017 10:46:10 PM	I liked a @YouTube vid youtu.be/xKjbLK52D1w'
4	galaxy s8 screen resolution confirmed in lates...	Link	844516071294074880	Original	Mar 22, 2017 7:48:16 AM	Galaxy S8 scre resolution confirme late:
5	мне понравилось видео "iphone 8 plus vs. galax...	No Media	912175638601195521	Original	Sep 25, 2017 12:43:14 AM	Мне понравилось вид "iPhone 8 Plus vs. Gala:
2	android circuit: new galaxy s8 issues, microso...	Image; Link	860637802731511808	Original	May 5, 2017 7:30:16 PM	Android Circuit: N Galaxy S8 Issu Microso
1	lito 3 in 1 electroplating hard pc phone case ...	Link	904644944434008066	Original	Sep 4, 2017 5:58:56 AM	LITO 3 in 1 Electroplati Hard PC Phone Case
8	i added a video to a @youtube playlist youtu.b...	Link	859054448336875520	Original	May 1, 2017 10:38:35 AM	I added a video to @YouTube play youtu.l
0	leak confirms iphone 8 will be larger than iph...	Link	870450876913156096	Original	Jun 1, 2017 9:23:55 PM	Leak Confirms iPhone Will Be Larger Than iPi
6	the new samsung galaxy s8 has finally arrived....	Image; Link	855416909285572609	Original	Apr 21, 2017 9:44:18 AM	The new Samsung Gal S8 has finally arrived
9	no root adblocker & package disabler - works o...	Link	915643244226269184	Original	Oct 4, 2017 2:22:15 PM	NO ROOT AdBlocke Package Disabler - wo (

10 rows × 23 columns

```
In [81]: for i in range(10):  
        text = random_tweets_df.iloc[i,5]  
        print(text)  
  
        for y in range(14,18):  
            print(random_tweets_df.columns[y],random_tweets_df.iloc[i,y])  
        print()
```

(10 iPhone HACKS and TRICKS 2017) has been published on My Iphone 8 - m  
yiphone8.co.uk/2017/10/15/10-... pic.twitter.com/xGGH3M88JX

TextBlob\_sentiment 0.0  
Afn\_Sent 0.0  
Vader\_Sentiment -0.3034  
Sentwordnet\_sent 0.0

My phone is beat to fuck...where this iPhone 8 at?

TextBlob\_sentiment 0.0  
Afn\_Sent -4.0  
Vader\_Sentiment 0.0  
Sentwordnet\_sent 0.06

California-Based Company Debuts New, Stylish Case Just in Time for the  
iPhone 8, iPhone 8 Plus and iPhone X Release ift.tt/2wZOIDF

TextBlob\_sentiment 0.3181818181818182  
Afn\_Sent 0.0  
Vader\_Sentiment 0.0  
Sentwordnet\_sent 0.07

Samsung Galaxy S8 SM-G950U - 64GB - Midnight Black (AT&T) Smartphone pi  
c.twitter.com/gvpDJKFzHC

TextBlob\_sentiment -0.16666666666666666  
Afn\_Sent 0.0  
Vader\_Sentiment 0.0  
Sentwordnet\_sent -0.07

Surprise: Galaxy S8 has the 'best smartphone display' buff.ly/2o2gJG1

TextBlob\_sentiment 1.0  
Afn\_Sent 3.0  
Vader\_Sentiment 0.743  
Sentwordnet\_sent 0.23

I liked a @YouTube video youtu.be/frbB28ofS\_c?a iPhone 8 Triple Bad New  
s Leak

TextBlob\_sentiment -0.049999999999999993  
Afn\_Sent -2.0  
Vader\_Sentiment -0.4767  
Sentwordnet\_sent -0.12

The LG V30 could get curves like the Galaxy S8 fb.me/8v8zRhmer

TextBlob\_sentiment 0.0  
Afn\_Sent 2.0  
Vader\_Sentiment 0.3612  
Sentwordnet\_sent 0.04

Galaxy S8 vs 7 Plus vs LG G6 vs Pixel vs 3T SPEED Test! youtu.be/OX4Juc  
pvpbJM via @YouTube Really great test of all the top phones.

TextBlob\_sentiment 0.65  
Afn\_Sent 5.0  
Vader\_Sentiment 0.7569  
Sentwordnet\_sent 0.09

Samsung's Galaxy S8 looks great but it still won't convince me to repla  
ce my iPhone. Maybe it's customer loyalty, maybe it's just right.

TextBlob\_sentiment 0.5428571428571429  
Afn\_Sent 10.0

Vader\_Sentiment 0.7343  
Sentwordnet\_sent -0.02

Apple unlikely to switch to USB-C on the iPhone 8 because you can't have nice things [dlvr.it/NWjGCG](http://dlvr.it/NWjGCG)  
TextBlob\_sentiment 0.04999999999999999  
Afn\_Sent 2.0  
Vader\_Sentiment 0.4215  
Sentwordnet\_sent 0.19

## 5. Apply Sentiment Analysis to Quality, Price, and Value of Devices

```
In [1]: import pandas as pd
all_data = pd.read_csv("all_data_sent.csv")
all_data.head()
```

C:\Users\dpc50\Anaconda3\lib\site-packages\IPython\core\interactiveshell.py:3020: DtypeWarning: Columns (7,9) have mixed types. Specify dtype option on import or set low\_memory=False.  
interactivity=interactivity, compiler=compiler, result=result)

Out[1]:

	LowerText	Media Type	Post ID	Post Type	Published Date (GMT-04:00) New York	Sound Bite Text	Source Type
0	following the naming system of the past severa...	Link	718bbba167877e763cfe851413849ed8	Original	May 8, 2017 7:39:37 AM	Following the naming system of the past severa...	Blogs
1	the processing cost for the oled-based 3d touc...	No Media	17194239754920322211	Original	May 19, 2017 7:42:09 AM	The processing cost for the OLED-based 3D Touc...	Blogs
2	the processing cost for the oled-based 3d touc...	No Media	5279493144373937172	Original	May 19, 2017 7:42:09 AM	The processing cost for the OLED-based 3D Touc...	Blogs
3	the processing cost for the oled-based 3d touc...	No Media	10937716975525293181	Original	May 19, 2017 7:46:00 AM	The processing cost for the OLED-based 3D Touc...	Blogs
4	we have 9 exciting iphone 8 rumors for the ult...	No Media	http://learnbonds.com/133761/iphone-8-rumors-f...	Original	May 20, 2017 10:20:00 AM	We Have 9 Exciting iPhone 8 Rumors for the Ult...	Blogs

```
In [2]: value = ["value", "worth", "use", "appreciate", "advantage", "benefit", "purpose"]
price = ["$", "price", "buy", "sell", "cost", "demand", "expensive", "cheap", "affordable", "money"]
quality = ["quality", "design", "display", "look", "screen", "battery", "water", "proof", "performance", "lens", "speaker"]
```

```
In [4]: samsung_data = all_data[all_data.device_type == "Samsung"]
```

```
In [6]: %%time
temp = None
for i, word in enumerate(price):
    if i == 0:
        temp = all_data.LowerText.str.contains(word)
    else:
        temp = temp | all_data.LowerText.str.contains(word)

all_data['price_text'] =temp
print(all_data['price_text'].sum())
```

91537

Wall time: 4.09 s

```
In [7]: %%time
temp = None
for i, word in enumerate(value):
    if i == 0:
        temp = all_data.LowerText.str.contains(word)
    else:
        temp = temp | all_data.LowerText.str.contains(word)

all_data['value_text'] =temp
print(all_data['value_text'].sum())
```

99398

Wall time: 2.73 s

```
In [8]: %%time
temp = None
for i, word in enumerate(quality):
    if i == 0:
        temp = all_data.LowerText.str.contains(word)
    else:
        temp = temp | all_data.LowerText.str.contains(word)

all_data['quality_text'] =temp
print(all_data['quality_text'].sum())
```

134186

Wall time: 4.31 s


```
In [9]: all_data.device_type.unique()
```

```
Out[9]: array(['Both iPhone 8 & iPhone X', 'Both Samsung & iPhone', 'iPhone 8',
              'iPhone X', 'Samsung'], dtype=object)
```

```
In [10]: import matplotlib.pyplot as plt
```


```
In [82]: all_data[all_data['price_text'] == True].TextBlob_sentiment.hist(bins = 20)
```

Out[82]: <matplotlib.axes.\_subplots.AxesSubplot at 0x249efc56860>



```
In [83]: all_data[all_data['value_text'] == True].TextBlob_sentiment.hist(bins = 20)
```


Out[83]: <matplotlib.axes.\_subplots.AxesSubplot at 0x249e6feec0>





```
In [84]: all_data[all_data['quality_text'] == True].TextBlob_sentiment.hist(bins
= 20)
```

```
Out[84]: <matplotlib.axes._subplots.AxesSubplot at 0x249f2abff28>
```



```
In [15]: all_data.columns
```

```
Out[15]: Index(['LowerText', 'Media Type', 'Post ID', 'Post Type',
'Published Date (GMT-04:00) New York', 'Sound Bite Text', 'Source
Type',
'contain8', 'containSamsung', 'containX', 'containiPhone',
'device_type', 'Tokenized_Text', 'Stemmed_text', 'TextBlob_senti
ment',
'Afn_Sent', 'Vader_Sentiment', 'Sentwordnet_sent', 'price_text',
'value_text', 'quality_text', 'Date'],
dtype='object')
```

## 6. Sentiment Analysis Before and After Releases

```
In [14]: %%time
all_data['Date'] = pd.to_datetime(all_data["Published Date (GMT-04:00) N
ew York"])
```

Wall time: 56.2 s

```
In [16]: samsung_release = "2017-04-21"
iphone_release = "2017-09-12"
```

```
In [92]: import numpy as np
samsung_data = []
features = ["price_text", "value_text", "quality_text"]
for feat in features:
    sub_data = all_data[all_data.device_type == "Samsung"]
    sub_data = sub_data[sub_data[feat] == True]
    samsung_data.append([sub_data[sub_data.Date < samsung_release].TextBlob_sentiment.mean(), sub_data[sub_data.Date > samsung_release].TextBlob_sentiment.mean()])


samsung_data = np.array(samsung_data)
```

```
In [93]: import numpy as np
iphone_data = []


for feat in features:
    sub_data = all_data[all_data.device_type == "iPhone 8"]
    sub_data = sub_data[sub_data[feat] == True]
    iphone_data.append([sub_data[sub_data.Date < iphone_release].TextBlob_sentiment.mean(), sub_data[sub_data.Date > iphone_release].TextBlob_sentiment.mean()])

iphone_data = np.array(iphone_data)
```

```
In [94]: n = np.array(list(range(3)))
plt.bar(n, samsung_data[:, 0], .35, label = "Pre-Release")
plt.bar(n+.35, samsung_data[:, 1], .35, label = "Post-Release")
plt.xticks(n+.35/2, features)
plt.legend()
plt.title("Samsung Galaxy S8 Sentiments Before and After Release")
plt.show()
```




```
In [95]: n = np.array(list(range(3)))
plt.bar(n,iphone_data[:,0],.35,label = "Pre-Release")
plt.bar(n+.35,iphone_data[:,1],.35,label = "Post-Release")
plt.xticks(n+.35/2,features)
plt.legend()
plt.title("iPhone 8 Sentiments Before and After Release")
plt.show()
```



```
In [96]: samsung_sent_pre_post = [all_data[all_data.device_type == "Samsung"][all_data[all_data.device_type == "Samsung"].Date < samsung_release].TextBlob_sentiment.mean(),all_data[all_data.device_type == "Samsung"][all_data[all_data.device_type == "Samsung"].Date > samsung_release].TextBlob_sentiment.mean()]
```

```
In [97]: iphone_8_sent_pre_post = [all_data[all_data.device_type == "iPhone 8"][all_data[all_data.device_type == "iPhone 8"].Date < iphone_release].TextBlob_sentiment.mean(),all_data[all_data.device_type == "iPhone 8"][all_data[all_data.device_type == "iPhone 8"].Date > iphone_release].TextBlob_sentiment.mean()]
```

```
In [138]: n = np.array(list(range(2)))
plt.bar(n,[samsung_sent_pre_post[0],iphone_8_sent_pre_post[0]],.35,label
= "Pre-Release")
plt.bar(n+.35,[samsung_sent_pre_post[1],iphone_8_sent_pre_post[1]],.35,l
abel = "Post-Release")
plt.xticks(n+.35/2,["Galaxy S8","iPhone 8"])
plt.legend()
plt.title("Galaxy S8 and iPhone 8 Sentiments Before and After Release")
plt.show()
```

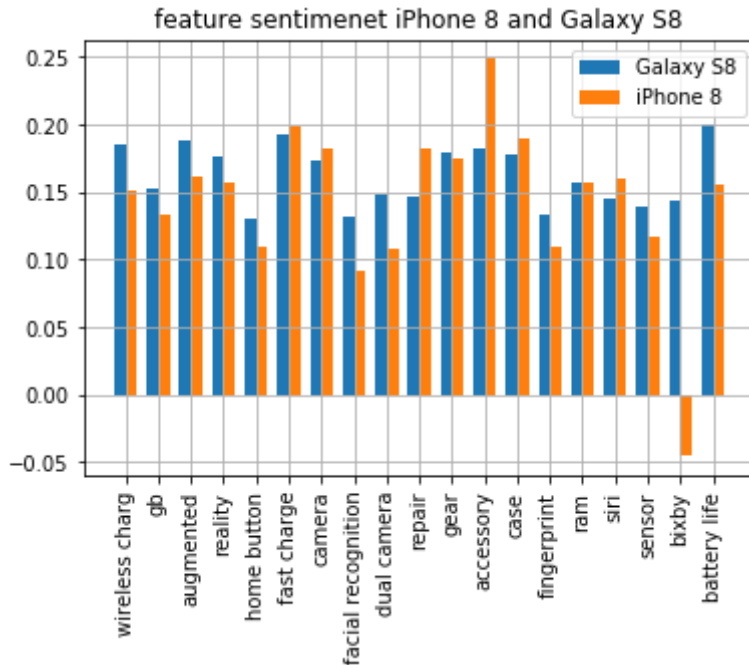


## 7. Sentiment of Features

```
In [101]: features = ["wireless charg","gb","augmented", "reality","home button",
"fast charge","camera","facial recognition","dual camera","repair","gea
r","accessory","case","fingerprint","ram","siri","sensor","bixby","batte
ry life"]
```

```
In [102]: feature_data = []
for f in features:
    row = []
    for dev in ["Samsung","iPhone 8"]:
        temp_df = all_data[all_data.device_type == dev]
        dev_val = temp_df[temp_df.LowerText.str.contains(f)].TextBlob_se
ntiment.mean()
        row.append(dev_val)
    feature_data.append(row)
feature_data = np.array(feature_data)
```

```
In [103]: width = .35
ind = np.arange(len(features))
plt.bar(ind,feature_data[:,0],width,label = "Galaxy S8")
plt.bar(ind+width,feature_data[:,1],width,label = "iPhone 8")
plt.xticks(ind+width/2,features,rotation = 90)
plt.title("feature sentiment iPhone 8 and Galaxy S8")
plt.legend()
plt.grid()
plt.show()
```




## 8. Twitter v. Non-Twitter data sources

```
In [104]: twitter_data = []
features = ["price_text", "value_text", "quality_text"]
for feat in features:
    sub_data = all_data[all_data["Source Type"] == "Twitter"]
    sub_data = sub_data[sub_data[feat] == True]
    twitter_data.append(sub_data.TextBlob_sentiment.mean())

twitter_data = np.array(twitter_data)
```

```
In [105]: non_twitter_data = []
features = ["price_text", "value_text", "quality_text"]
for feat in features:
    sub_data = all_data[all_data["Source Type"] != "Twitter"]
    sub_data = sub_data[sub_data[feat] == True]
    non_twitter_data.append(sub_data.TextBlob_sentiment.mean())
non_twitter_data = np.array(non_twitter_data)
```

```
In [106]: n = np.array(list(range(3)))
plt.bar(n,twitter_data,.35,label = "Twitter")
plt.bar(n+.35,non_twiter_data,.35,label = "Non-Twitter")
plt.xticks(n+.35/2,["Price","Value","Quality"])
plt.legend()
plt.title("Twitter v Non-Twtter Sentiments")
plt.show()
```



## 9. Word Cloud Data Exploration

```
In [107]: from wordcloud import WordCloud, STOPWORDS
stopwords = list(STOPWORDS) + ["iphone", "galaxy", "s8", "ifttt", "http", "if
t", "tt", "samsung"]
def wordcloud(text,col,stopwords):
    wordcloud = WordCloud(background_color="white",stopwords=stopwords).
generate(" ".join([i for i in text[col]]))
    plt.figure( figsize=(20,10), facecolor='k')
    plt.imshow(wordcloud)
    plt.axis("off")
```


```
In [108]: for device in all_data.device_type.unique():  
          temp = all_data[all_data.device_type == device]  
          print(device)  
          wordcloud(temp, "LowerText", stopwords)  
          plt.show()
```

[illegible][illegible]

25/32







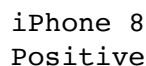



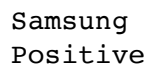
```
In [110]: #positive and negative for all devices
for device in all_data.device_type.unique():
    temp = all_data[all_data.device_type == device]
    if device != "iPhone X":
        print(device)
        print("Positive")
        positive_temp = temp[temp.TextBlob_sentiment > temp.TextBlob_sentiment.mean()+ 2*temp.TextBlob_sentiment.std()]
        wordcloud(positive_temp, "LowerText", stopwords)

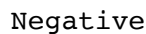
        plt.show()
        print("Negative")
        neg_temp = temp[temp.TextBlob_sentiment < temp.TextBlob_sentiment.mean()- 2*temp.TextBlob_sentiment.std()]
        wordcloud(neg_temp, "LowerText", stopwords)

        plt.show()
```









# 95-851: Social Media Analysis

## Data Clan

### Part 2: Important Attribute Analysis

## 1. Load Data

```
In [0]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import spacy
import seaborn as sns
```

```
In [0]: import re
import string
import codecs

def makeWordList(path):

    with codecs.open(path, "r", encoding='utf-8', errors='ignore') as f:
        corpus_text = f.read()

    for c in string.punctuation:
        corpus_text = corpus_text.replace(c, "") # -- (1)

    text = re.sub(r'\S*\d\S*', '', corpus_text) # -- (2)
    text = re.sub(r'^\w\s', '', text) # -- (3)

    text = text.lower().split() # -- (4)

    li = []
    for token in text:
        li.append(token)

    return " ".join(li)
```

```
In [0]: # load text corpus
#Both_iPhones_data
df = pd.read_csv('./Cleaned/Both_iPhones_data.csv', encoding='utf-8')
```


## 2. iPhone Analysis



## Sources of iPhone posts:

```
In [0]: df['Source Type'].value_counts().plot('bar')
```

```
Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7f8af6057128>
```



As we can see majority of the reviews/posts come from blogs.

```
In [0]: corpus = list(df['Sound Bite Text'])
```

```
In [0]: corpus[0]
```

```
Out[0]: 'Following the naming system of the past several years of the iPhone, this year's release would be named the iPhone 7s. Talk suggests there will still be the release of the iPhone 7s in the addition to the newer model. What this will be named – iPhone 8, iPhone 10, iPhone X, something completely different – we'll just have to wait and see! Lesson of the Week: Julian Robertson "Our mandate is to find the 200 best companies in the world and invest in them, and find the 200 worst companies in the world and go short on them.'
```

```
In [0]: len(corpus)
```

```
Out[0]: 38405
```

There 38405 reviews/posts on iPhone 8 and iPhone X combined

```
In [0]: import spacy
nlp = spacy.load('en', disable=['ner', 'parser', 'tagger']) #disabling t
o optimize loading time

def word_tokenizer(doc):
    parsed_doc = nlp(doc)
    return([token.lemma_.lower() for token in parsed_doc if re.match('[a
-zA-Z]+$', token.orth_) and token.lemma_ != '-PRON-'])
```

```
In [0]: from sklearn.feature_extraction import text
my_stop_words = text.ENGLISH_STOP_WORDS.union(['iphone', 'apple', 'mac',
'ipad', 'plus', 'iphonex', 'applewatch', 'macbook'])
```

```
In [0]: #constructing a document matrix
from sklearn.feature_extraction.text import TfidfVectorizer
vectorizer = TfidfVectorizer(min_df=50, stop_words=my_stop_words, max_df
=0.8, tokenizer=word_tokenizer)
doc_matrix = vectorizer.fit_transform(corpus)
print('Number of features in tf-idf:', len(vectorizer.vocabulary_)) # a
mapping of terms to feature indices
```

Number of features in tf-idf: 3324

```
In [0]: doc_matrix = doc_matrix.toarray()
```

```
In [0]: doc_matrix.shape
```

```
Out[0]: (38405, 3324)
```

After feature engineering, we have extract 3324 feature values corresponding to 38405 posts/reviews.

```
In [0]: #constructing feature set
feature_set = pd.DataFrame(doc_matrix)
```

```
In [0]: feature_set.columns = vectorizer.get_feature_names()
```

```
In [0]: feature_set.head()
```

```
Out[0]:
```

	10	2	3	4	5	a.m.	aapl	abandon	ability	able	...	youtuber	yuan	zaharov
0	0.000000	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
1	0.085010	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
2	0.091318	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
3	0.090216	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
4	0.000000	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0

5 rows × 3324 columns

```
In [0]: num_topics = 10
        from sklearn.decomposition import LatentDirichletAllocation
        lda = LatentDirichletAllocation(n_components=num_topics, learning_method
        ='online', max_iter=10, n_jobs=-1, random_state=0)
        lda.fit(doc_matrix)
```

```
Out[0]: LatentDirichletAllocation(batch_size=128, doc_topic_prior=None,
        evaluate_every=-1, learning_decay=0.7,
        learning_method='online', learning_offset=10.0,
        max_doc_update_iter=100, max_iter=10, mean_change_tol=0.00
        1,
        n_components=10, n_jobs=-1, n_topics=None, perp_tol=0.1,
        random_state=0, topic_word_prior=None,
        total_samples=1000000.0, verbose=0)
```

```
In [0]: print("Log Likelihood: ", lda.score(doc_matrix))

Log Likelihood: -1577041.7641789156
```

```
In [0]: print("Perplexity: ", lda.perplexity(doc_matrix))

Perplexity: 3110.62261221589
```

```
In [0]: best_topics = 10
        best_lda_model = lda
```

```
In [0]: # Topic - Word matrix
        topic_keywords_matrix = pd.DataFrame(best_lda_model.components_)
        topic_keywords_matrix.columns = vectorizer.get_feature_names()
```

```
In [0]: topic_word_distributions = np.array([topic_word_pseudocounts / np.sum(to
        pic_word_pseudocounts)
                                              for topic_word_pseudocounts in best
        _lda_model.components_])
```

```
In [0]: word_columns = vectorizer.get_feature_names()
```

## Classify posts based on common themes

```
In [0]: # Top 20 words per topic
print('Displaying the top %d words per topic and their probabilities within the topic...' % 20)
print()

for topic_idx in range(best_topics):
    print('[Topic ', topic_idx, ']', sep='')
    sort_indices = np.argsort(topic_word_distributions[topic_idx][:,-1])
    for rank in range(20):
        word_idx = sort_indices[rank]
        print(word_columns[word_idx], ':', topic_word_distributions[topic_idx, word_idx])
    print()
```

Displaying the top 20 words per topic and their probabilities within the topic...

[Topic 0]

gb : 0.08124314795349012  
gold : 0.04410283249794682  
silver : 0.03711345405706172  
iphonex : 0.0285691050166316  
grey : 0.02727555025725522  
ready : 0.026801837573345554  
warranty : 0.021902380389274053  
idr : 0.02064325040860555  
space : 0.018827639424979174  
black : 0.01584397212368294  
color : 0.01509612515748844  
appletv : 0.013902717245887595  
po : 0.013067984098703432  
plus : 0.011254466475556044  
instagood : 0.010192217994243068  
order : 0.010189852782789836  
eta : 0.0098416074187699  
available : 0.009334004021601942  
open : 0.009299183846883606  
jualiphone : 0.008682850861908217

[Topic 1]

apple : 0.01815531268165605  
new : 0.0086841931537251  
plus : 0.008435350915833199  
price : 0.008331457723834815  
launch : 0.008097981188887608  
order : 0.00793602440560132  
pre : 0.0067097997776872885  
buy : 0.00602034729244769  
release : 0.005694689001243799  
say : 0.005668815577593309  
store : 0.005218897977990837  
year : 0.0051913922843324556  
model : 0.0050264809302259204  
report : 0.004999077140273415  
phone : 0.004960685616460343  
start : 0.0046675662288371556  
watch : 0.004647171762996064  
demand : 0.004630651588159105  
month : 0.0046294769896978736  
cost : 0.00461872773267132

[Topic 2]

apple : 0.011828223378695959  
new : 0.010006526162587823  
plus : 0.009257686725840368  
phone : 0.00669340601564742  
camera : 0.006646041024878505  
feature : 0.005944818032704595  
screen : 0.005280006789421025  
make : 0.0046961159768002724  
use : 0.004618115398280779

device : 0.004578096433708312  
ios : 0.004558068432185897  
display : 0.0043804375700023535  
like : 0.004282907722111291  
design : 0.00420636343799351  
face : 0.004029503978618222  
look : 0.0037913617336497673  
come : 0.003764873922151073  
pixel : 0.003695638816508181  
just : 0.00366095659509803  
charge : 0.0036340945592757924

[Topic 3]

charge : 0.09471580298408803  
wireless : 0.08809183515390136  
qi : 0.03385878645396637  
pad : 0.03366664704402715  
charger : 0.02851052747205663  
apple : 0.020380056723358932  
powerbyproxi : 0.019116451497411972  
support : 0.018801019730501457  
ikea : 0.01604313863402815  
case : 0.01563248848028313  
standard : 0.015162903208354289  
include : 0.01332835582450516  
new : 0.011165936727383573  
accessory : 0.011103829175289681  
belkin : 0.011011322946472989  
airpower : 0.01099039025909077  
plug : 0.010038895446879493  
power : 0.009287162741726637  
work : 0.008985304164221561  
starbucks : 0.008754326118411045

[Topic 4]

iphonex : 0.08960288972084668  
plus : 0.04746456270083673  
tag : 0.038212148016760615  
apple : 0.033074284759226996  
vs : 0.028523997028196503  
ifttt : 0.02745385169595188  
twitter : 0.019510303088625374  
tech : 0.019205887195247066  
ios : 0.01815838823625497  
macbook : 0.017664221001397994  
unboxing : 0.01727397601811946  
beat : 0.01686242554196894  
applewatch : 0.015469643578818509  
buy : 0.015277808763483421  
ipad : 0.014220181532959767  
video : 0.014144355018765547  
technology : 0.01341462433820431  
red : 0.013404207616320099  
fee : 0.013206740931377792  
photo : 0.012973578849661068

[Topic 5]

charge : 0.044324826135869765  
usb : 0.04097267300820534  
fast : 0.039901188917655005  
c : 0.03329894061261873  
charger : 0.02686114602812956  
adapter : 0.02356626362105339  
cable : 0.01936804869551776  
power : 0.01798376715485782  
keybanc : 0.016901642051168744  
lightning : 0.016680226841080698  
support : 0.014710136949003862  
minute : 0.014507291653038294  
post : 0.014431883955452184  
plus : 0.014370831190036158  
review : 0.014322574861014894  
l : 0.013507167036357503  
apple : 0.01321933361759763  
box : 0.012828378503101917  
need : 0.012468429764443343  
overload : 0.012327821276818402

#### [Topic 6]

irrelevant : 0.043299476684149324  
flag : 0.04311701274629856  
google : 0.03838700904319528  
alert : 0.030528258504849335  
news : 0.026788452297191675  
cnet : 0.022953801970168143  
ifttt : 0.020759165189906236  
business : 0.018242310135475937  
mobile : 0.01608912860332735  
coverage : 0.015969745398043706  
forbes : 0.015771801605714935  
tag : 0.015555109796640867  
t : 0.015299204936198288  
insider : 0.014314094546552858  
network : 0.012905734364261397  
apple : 0.01275583378990395  
lte : 0.01270380479182214  
blog : 0.010880641166205143  
app : 0.010836862845147132  
support : 0.010454659010314707

#### [Topic 7]

survey : 0.0488375331585272  
bio : 0.030879927921308923  
link : 0.023593061094276523  
decline : 0.02272482765930906  
poll : 0.016621392092961748  
percent : 0.016245034958426538  
life : 0.01618450352169856  
pick : 0.016011803060931615  
profit : 0.015032429903392586  
euro : 0.014788897081045402  
guy : 0.014353902788908118  
bernstein : 0.013664133065030348  
engadget : 0.01352622911061442

fm : 0.013343328720601282  
battery : 0.012951714084298236  
reaction : 0.012914364971852582  
quick : 0.01284860636083979  
want : 0.012751589472656784  
leader : 0.012216763718522117  
capacity : 0.012203182161546963

[Topic 8]

stock : 0.03585405365791643  
finance : 0.0345413882921526  
white : 0.027059329899661453  
dan : 0.025603750948926175  
q : 0.02313041247333707  
complicate : 0.020661015560073832  
trust : 0.020142604216971095  
ini : 0.019370501565060195  
quote : 0.01898964986396941  
wife : 0.01716157180700678  
crackle : 0.016350318296938463  
murata : 0.01607455274165567  
supplier : 0.01571600155834047  
arab : 0.014150862974402172  
limit : 0.013656971443163024  
anda : 0.012738443695236875  
hari : 0.012577921305126055  
kali : 0.012394704461044157  
jackpot : 0.01219542364602757  
doom : 0.011781556661991209

[Topic 9]

lol : 0.03301548070797941  
nintendo : 0.023786806892505313  
love : 0.0235321469320083  
wait : 0.01920062996033466  
upsell : 0.018962421686790593  
okami : 0.0187734724978438  
patent : 0.018483969696744904  
pink : 0.01834346037205404  
hardwick : 0.018191343816844657  
rating : 0.017932316790226475  
lesson : 0.01735694558175868  
silicon : 0.016374992164142396  
story : 0.016045993431963248  
valley : 0.01579357132723258  
syndicate : 0.015691816928065185  
tim : 0.014261446137632526  
promo : 0.013799462239500314  
viral : 0.013332275540478089  
original : 0.012487001474789717  
xs : 0.012091667016089382



Topic 0 gives us an idea that users are talking about color of the phone here. Top keywords are 'black' 'color', 'gold' color, 'silver' color, 'warranty' and 'space'.

Topic 1 tells us that here users are talking about price, cost etc.

Topic 2 tells us that here users are talking about display, design, look, pixel, screen etc.

Topic 3 tells us that users here are talking about charge/charging, plug, power, wireless

Topic 4 seems irrelevant to our analysis

Topic 5 users are talking about charging/power similar to topic 3.

Topic 6, 7, 8, 9 seems irrelevant.

### Experimentation with bigrams (instead of unigrams as per above)

```
In [0]: from sklearn.feature_extraction.text import CountVectorizer
countvectorizer = CountVectorizer(min_df=50, stop_words=my_stop_words, m
ax_df=0.8, tokenizer=word_tokenizer, ngram_range=(2,2))
count_matrix = countvectorizer.fit_transform(corpus)
print('Number of features in tf-idf:', len(countvectorizer.vocabulary_))
# a mapping of terms to feature indices
```

Number of features in tf-idf: 4668

```
In [0]: count_matrix = count_matrix.toarray()
sum_bigrams = count_matrix.sum(axis=0)
words_freq = [(word, sum_bigrams[idx]) for word, idx in countvectorizer.
vocabulary_.items()]
words_freq =sorted(words_freq, key = lambda x: x[1], reverse=True)
```

```
In [0]: words_freq[:30]
```

```
Out[0]: [('wireless charge', 10643),
 ('pre order', 6228),
 ('new x', 3829),
 ('watch series', 3683),
 ('new iphones', 3570),
 ('gb gb', 3536),
 ('home button', 2675),
 ('augment reality', 2482),
 ('new phone', 2421),
 ('new feature', 2406),
 ('fast charge', 2285),
 ('flag irrelevant', 2185),
 ('portrait light', 1944),
 ('edge edge', 1775),
 ('appear \ufe01', 1742),
 ('bionic chip', 1730),
 ('portrait mode', 1699),
 ('facial recognition', 1689),
 ('battery life', 1681),
 ('charge pad', 1678),
 ('steve job', 1613),
 ('new model', 1507),
 ('dual camera', 1455),
 ('x come', 1422),
 ('wireless charger', 1371),
 ('usb c', 1365),
 ('x x', 1335),
 ('tim cook', 1290),
 ('oled display', 1278),
 ('brand new', 1278)]
```

## Most important attributes for iPhone:

The 30 most frequently occurring bigrams in the reviews/post based on the above analysis. This gives us a hint on what features/characteristics of the phone are being discussed by the users the most. Some of the most talked about attributes based on our analysis above were:

1. Wireless charging
2. Home button (Apple removed home button from iPhone X, so this generated a good amount of discussion among users)
3. Fast Charge
4. Portrait lighting (New feature introduced in iPhone)
5. Bionic Chip (Apple introduced this new feature iPhone 8 onwards)
6. Facial recognition (Apple introduced this new feature iPhone 8 onwards)
7. Charge Pad (Apple introduced this new feature iPhone 8 onwards)
8. Dual Camera
9. Oled Display (New OLED screen for iPhone X)

## Perform bigram topic modeling

```
In [0]: #constructing a document matrix
vectorizer2 = TfidfVectorizer(min_df=50, stop_words=my_stop_words, max_d
f=0.8, tokenizer=word_tokenizer, ngram_range=(2,2))
doc_matrix2 = vectorizer2.fit_transform(corpus)
print('Number of features in tf-idf:', len(vectorizer2.vocabulary_)) # a
mapping of terms to feature indices
```

Number of features in tf-idf: 4668

```
In [0]: doc_matrix2 = doc_matrix2.toarray()
```

```
In [0]: doc_matrix2.shape
```

Out[0]: (38405, 4668)

After feature engineering, we have extract 3324 feature values corresponding to 38405 posts/reviews.

```
In [0]: #constructing feature set
feature_set2 = pd.DataFrame(doc_matrix2)
```

```
In [0]: feature_set2.columns = vectorizer2.get_feature_names()
```

```
In [0]: feature_set2.head()
```

Out[0]:

	10 anniversary	2 generation	2 new	3 generation	3 party	3 quarter	4 quarter	aapl stock	able charge	able make	...	n
0	0.000000	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	
1	0.133166	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	
2	0.136158	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	
3	0.136158	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	
4	0.000000	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	

5 rows × 4668 columns

```
In [0]: num_topics = 10
        from sklearn.decomposition import LatentDirichletAllocation
        lda2 = LatentDirichletAllocation(n_components=num_topics, learning_method='online', max_iter=10, n_jobs=-1, random_state=0)
        lda2.fit(doc_matrix2)
```

```
Out[0]: LatentDirichletAllocation(batch_size=128, doc_topic_prior=None,
                                   evaluate_every=-1, learning_decay=0.7,
                                   learning_method='online', learning_offset=10.0,
                                   max_doc_update_iter=100, max_iter=10, mean_change_tol=0.00
1,
                                   n_components=10, n_jobs=-1, n_topics=None, perp_tol=0.1,
                                   random_state=0, topic_word_prior=None,
                                   total_samples=1000000.0, verbose=0)
```

```
In [0]: print("Log Likelihood: ", lda2.score(doc_matrix2))

Log Likelihood:  -1118189.9989448085
```

```
In [0]: print("Perplexity: ", lda2.perplexity(doc_matrix2))

Perplexity:  8249.049629477855
```

```
In [0]: # Select the best hyperparameters
        best_topics2 = 10
        best_lda_model2 = lda2
```

```
In [0]: # Topic - Word matrix
        topic_keywords_matrix2 = pd.DataFrame(best_lda_model2.components_)
        topic_keywords_matrix2.columns = vectorizer2.get_feature_names()
```

```
In [0]: topic_word_distributions2 = np.array([topic_word_pseudocounts / np.sum(topic_word_pseudocounts)
                                                for topic_word_pseudocounts in best_lda_model2.components_])
```

```
In [0]: word_columns2 = vectorizer2.get_feature_names()
```

```
In [0]: # Top 20 words per topic
print('Displaying the top %d bigrams per topic and their probabilities within the topic...' % 20)
print()

for topic_idx in range(best_topics):
    print('[Topic ', topic_idx, ']', sep='')
    sort_indices2 = np.argsort(topic_word_distributions2[topic_idx][::-1])
    for rank in range(20):
        word_idx = sort_indices2[rank]
        print(word_columns2[word_idx], ': ', topic_word_distributions2[topic_idx, word_idx])
    print()
```

Displaying the top 20 bigrams per topic and their probabilities within the topic...

[Topic 0]

wireless charge : 0.01777715387585753  
fast charge : 0.005707561620969423  
home button : 0.005619200957530918  
new x : 0.005398372501192109  
new iphones : 0.00480258993425143  
new feature : 0.004261652655317761  
watch series : 0.004260063519744209  
battery life : 0.004204906541827839  
charge pad : 0.004000343630079443  
bionic chip : 0.0037819184737335715  
oled display : 0.0037052513743764666  
facial recognition : 0.0033390255074162135  
edge edge : 0.003289685871626061  
dual camera : 0.003283278056528878  
usb c : 0.003269607419867699  
qi wireless : 0.0031986219364717925  
new phone : 0.0031846874940162266  
x screen : 0.0031141966471887146  
x new : 0.0030215366906442506  
augment reality : 0.002967439346028165

[Topic 1]

gb gb : 0.03762385944149371  
t mobile : 0.01874271496916777  
space grey : 0.01840246092550391  
click expand : 0.01592886398264929  
reason buy : 0.015287086148785026  
press release : 0.012144052130992962  
x tag : 0.01186034899489786  
make sense : 0.011428131604306014  
gold silver : 0.011185669310730267  
gb x : 0.009612922557250797  
buy instead : 0.009541162214916148  
gb idr : 0.009488791872972623  
order x : 0.009434609995557806  
x gb : 0.00930293733367522  
grey gb : 0.008380105860010175  
early month : 0.008249423145602989  
silver black : 0.008136187753191946  
release blog : 0.00795947806393744  
support new : 0.007949153367450409  
x compare : 0.007947430379486054

[Topic 2]

view wordpress : 0.01516915329326385  
wordpress tag : 0.011142057294053076  
x cost : 0.011084407682524864  
repair cost : 0.010461507821954702  
steve job : 0.010063961765172498  
story watch : 0.009969386694481205  
conversation story : 0.009621811321003268  
join conversation : 0.009603545108181538  
x important : 0.009017556228525967

screen repair : 0.00891184053788994  
x device : 0.007833158801703564  
watch announce : 0.0077626973491623135  
plan buy : 0.007459852475847686  
cost repair : 0.007398801985339411  
tech insider : 0.007206758638818512  
like yes : 0.006364975277894009  
job theater : 0.006302653246049779  
social medium : 0.006039151765313626  
double wireless : 0.005870391854396006  
tech gadget : 0.005470332624200753

#### [Topic 3]

pre order : 0.0367942741361525  
x vs : 0.018229173819462798  
x pre : 0.013428087963948553  
buy x : 0.012697034576084428  
x launch : 0.010380085337237341  
launch day : 0.008403785174313002  
vs x : 0.008359200815876662  
release date : 0.007577479837023744  
include x : 0.007470432795633398  
ios device : 0.007412353098381237  
x available : 0.006939876196166203  
brand new : 0.006884312128899353  
appear 1 : 0.0067936578660800365  
red picture : 0.0064154740272507815  
ahead x : 0.0061044830376420145  
people wait : 0.00606742215131575  
new x : 0.006065973307130645  
available pre : 0.005905824716179976  
x production : 0.005744568440631799  
new ios : 0.005712308997467998

#### [Topic 4]

watch series : 0.009976390732677577  
pre order : 0.008424344894718007  
pixel xl : 0.00796612751089979  
launch x : 0.006743084886250254  
x watch : 0.006360760923148467  
x price : 0.006017138045293169  
announce new : 0.00597360980508161  
release x : 0.005720383445273442  
new x : 0.00547676867527389  
phone x : 0.0052433276274147286  
x release : 0.005082340676331869  
gb model : 0.005066371289838914  
x x : 0.004870949980700288  
wait x : 0.004864324170263605  
wireless charger : 0.004431993943319008  
tag ifttt : 0.004392754196807779  
new model : 0.004354409017870886  
announce x : 0.004344413471617086  
unveil new : 0.004324292015832824  
upcoming x : 0.004256955761630575

#### [Topic 5]

wait x : 0.022631219614314073  
google news : 0.022206278821022774  
tag ifttt : 0.020434267608786565  
technology google : 0.02042808135488185  
google pixel : 0.017994727620403272  
ifttt tag : 0.01598691155204227  
fee ifttt : 0.012294987609380881  
news article : 0.012094114338428674  
x look : 0.010901292967946811  
coverage technology : 0.010813969147249313  
tear x : 0.010505070051334308  
upgrade x : 0.009685056333560317  
force restart : 0.009138265139555348  
compare x : 0.008793762135525773  
x follow : 0.008716006885858196  
article technology : 0.00823009142113509  
wi fi : 0.008119325045843367  
late iphones : 0.00793082474590451  
article forum : 0.007159186203082063  
x vs : 0.007005076857336505

#### [Topic 6]

new x : 0.01845351207275172  
x early : 0.014661865895045128  
x tech : 0.012515804356585222  
case x : 0.010324525609612646  
tech news : 0.01027652703416699  
read original : 0.009713569900372494  
x read : 0.009384985023621178  
x begin : 0.009056008483964591  
original post : 0.008344906105505649  
want new : 0.008298202891983481  
await x : 0.008207057398361929  
x actually : 0.008069475796140722  
x announcement : 0.008034323178426266  
long await : 0.007749218447140597  
launch month : 0.007376135311276976  
thing x : 0.007237513006818029  
make sure : 0.007152039470720113  
x ready : 0.007126510222274577  
announcement x : 0.006864795729980973  
swell battery : 0.006629691124221259

#### [Topic 7]

tag x : 0.018665099637521994  
tech technology : 0.014550602618738065  
want x : 0.012767396766025646  
airpods appletv : 0.010861281545403903  
appletv red : 0.010556565149701215  
beat photo : 0.010556565149680845  
picture beat : 0.010556565149680642  
ios airpods : 0.010556565149671031  
publish 1 : 0.008542896231318486  
x event : 0.008444574976048287  
fan hold : 0.008221432928863242  
information overload : 0.007783837074654198  
overload news : 0.007783837074654198



x think : 0.006938942437503515  
smartphone x : 0.006747806933323833  
tag tech : 0.006542870096404709  
rss fee : 0.006514280613442105  
protective case : 0.006307453756539623  
buy wait : 0.0062063234859421855  
receive email : 0.005926756019365995

[Topic 8]

flag irrelevant : 0.030978940699707518  
augment reality : 0.011661082738129349  
tim cook : 0.01136735869853837  
link bio : 0.010229944460782863  
ming chi : 0.009728043494038647  
chi kuo : 0.00970624856372086  
supply chain : 0.008969505134708233  
demand x : 0.00863925299938879  
post x : 0.0072906979263332485  
cut production : 0.007203029888510839  
analyst ming : 0.0070870885803032  
x announce : 0.0068646699694821815  
kgi security : 0.006765120904902963  
million unit : 0.006365353063468234  
video x : 0.006239646971532049  
capital market : 0.006202412541195928  
coverage flag : 0.006097327642347377  
x edition : 0.0059349168667481795  
x unit : 0.005693431110674264  
high price : 0.005580698483380989

[Topic 9]

portrait light : 0.02052770099916783  
x case : 0.019758721601113926  
gb ram : 0.013363843453209106  
continue read : 0.009734371336432157  
low light : 0.009350199449015556  
x come : 0.008806243167385213  
case available : 0.008392627566891964  
like x : 0.008135730268508837  
car support : 0.007884930348044307  
light effect : 0.007086515472157515  
phone like : 0.007041962666465463  
new video : 0.006985850627826908  
portrait mode : 0.006970989617306466  
x include : 0.006427039202337871  
new pixel : 0.0063420966416683645  
screen protector : 0.006308424161221779  
allow user : 0.00605734480695968  
temper glass : 0.005966556480412648  
light mode : 0.005964998936192403  
power button : 0.005724769985985333


## 2. Galaxy S8 Analysis

```
In [0]: # load text corpus
# samsung_data
df = pd.read_csv('./Cleaned/samsung_data.csv', encoding='utf-8')
```

### Sources of Galaxy S8 posts:

```
In [0]: df['Source Type'].value_counts().plot('bar')
```

```
Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7f8a1d57fd0>
```



As we can see majority of the reviews/posts come from blogs.

```
In [0]: corpus = list(df['Sound Bite Text'])
```

```
In [0]: corpus[0]
```

```
Out[0]: 'Samsung Galaxy S8 G950U G950P Sprint unlock by usb cable On Cell Phone
Forums Need to unlock a S8 Sprint ? Find here an exclusive Sprint Samsu
ng Galaxy S8 G950U G950P unlock method, available for this phone model
that cannot be unlocked by unlock codes. This Samsung Galaxy S8 unlock
service is the only option to unlock this phone model and make it work
both in the USA and abroad on any gsm network. The SM-G950U G950P unloc
k is done online by USB cable, not by codes calculated from IMEI. ... S
upported models - Sprint Samsung Galaxy S8 G950U G950P HOW THE SAMSUNG
GALAXY S8 UNLOCK IS DONE? Important: We are available according to our
Online program , there is shown if are available at a certain time. Aft
er buying this Sprint Samsung s8 unlock service, you will receive in a
few minutes the instructions about how to prepare your computer and pho
ne before the online unlock (some simple steps in our video tutorial, t
akes less than 5 min).'
```

```
In [0]: len(corpus)
```

```
Out[0]: 159419
```

There 159419 reviews/posts on Samsung

```
In [0]: import spacy
nlp = spacy.load('en', disable=['ner', 'parser', 'tagger']) #disabling t
o optimize loading time

def word_tokenizer(doc):
    parsed_doc = nlp(doc)
    return([token.lemma_.lower() for token in parsed_doc if re.match('[a
-zA-Z]+$', token.orth_) and token.lemma_ != '-PRON-'])
```

```
In [0]: from sklearn.feature_extraction import text
my_stop_words = text.ENGLISH_STOP_WORDS.union(['android', 'samsung', 'ga
laxy'])
```

```
In [0]: #constructing a document matrix
from sklearn.feature_extraction.text import TfidfVectorizer
vectorizer = TfidfVectorizer(min_df=50, stop_words=my_stop_words, max_df
=0.8, tokenizer=word_tokenizer, max_features=10000)
doc_matrix = vectorizer.fit_transform(corpus)
print('Number of features in tf-idf:', len(vectorizer.vocabulary_)) # a
mapping of terms to feature indices
```

```
Number of features in tf-idf: 5893
```

```
In [0]: doc_matrix = doc_matrix.toarray()
```

```
In [0]: doc_matrix.shape
```

```
Out[0]: (159419, 5893)
```

After feature engineering, we have extract 3324 feature values corresponding to 38405 posts/reviews.

```
In [0]: #constructing feature set
feature_set = pd.DataFrame(doc_matrix)
```

```
In [0]: feature_set.columns = vectorizer.get_feature_names()
```

```
In [0]: feature_set.head()
```

```
Out[0]:
```

	2	3	4	5	8	a.m.	abandon	abhijeet	ability	able	...	zenfone	zenpad	zero	zip
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0

5 rows × 5893 columns

```
In [0]: num_topics = 10
from sklearn.decomposition import LatentDirichletAllocation
lda = LatentDirichletAllocation(n_components=num_topics, learning_method='online', max_iter=10, n_jobs=-1, random_state=0)
lda.fit(doc_matrix)
```

```
Out[0]: LatentDirichletAllocation(batch_size=128, doc_topic_prior=None,
    evaluate_every=-1, learning_decay=0.7,
    learning_method='online', learning_offset=10.0,
    max_doc_update_iter=100, max_iter=10, mean_change_tol=0.00
1,
    n_components=10, n_jobs=-1, n_topics=None, perp_tol=0.1,
    random_state=0, topic_word_prior=None,
    total_samples=1000000.0, verbose=0)
```

```
In [0]: print("Log Likelihood: ", lda.score(doc_matrix))
```

Log Likelihood: -5798516.057844428

```
In [0]: print("Perplexity: ", lda.perplexity(doc_matrix))
```

Perplexity: 4538.315328002736

```
In [0]: # Select the best hyperparameters
best_topics = 10
best_lda_model = lda
```

```
In [0]: # Topic - Word matrix
topic_keywords_matrix = pd.DataFrame(best_lda_model.components_)
topic_keywords_matrix.columns = vectorizer.get_feature_names()
```

```
In [0]: topic_word_distributions = np.array([topic_word_pseudocounts / np.sum(topic_word_pseudocounts)
    for topic_word_pseudocounts in best_lda_model.components_])
```

```
In [0]: word_columns = vectorizer.get_feature_names()
```

## **Classify posts based on common themes**

```
In [0]: # Top 20 words per topic
print('Displaying the top %d words per topic and their probabilities within the topic...' % 20)
print()

for topic_idx in range(best_topics):
    print('[Topic ', topic_idx, ']', sep='')
    sort_indices = np.argsort(topic_word_distributions[topic_idx][:,-1])
    for rank in range(20):
        word_idx = sort_indices[rank]
        print(word_columns[word_idx], ':', topic_word_distributions[topic_idx, word_idx])
    print()
```

Displaying the top 20 words per topic and their probabilities within the topic...

[Topic 0]

order : 0.025629200779942435  
pre : 0.023921682668846124  
red : 0.0200222946339203  
update : 0.019982828284417258  
tint : 0.017059332137557472  
mobile : 0.016332604114190775  
fix : 0.014775292640924982  
issue : 0.014586623602296258  
t : 0.014016315097136106  
april : 0.012957829245554362  
unlock : 0.011438362430679635  
carrier : 0.011053283721860265  
available : 0.010391017561181994  
verizon : 0.00990162520947982  
preorder : 0.009302892980303549  
start : 0.009047962760063819  
plus : 0.008824077481106484  
free : 0.007711189121191785  
buy : 0.007628000092502139  
customer : 0.007606230136410973

[Topic 1]

dex : 0.05082160971296733  
desktop : 0.030507896922510797  
dock : 0.028539222046489263  
cnet : 0.026972042363368698  
send : 0.024042807306729905  
icon : 0.023636572949141302  
launcher : 0.022831299760820677  
use : 0.022076469307822748  
station : 0.02028452643008502  
gagt : 0.01979355103667749  
pc : 0.012929465571382692  
theme : 0.012714763966309647  
pic : 0.011716710352179905  
oneplus : 0.011575105289503715  
blow : 0.01039817042052308  
keyboard : 0.009945257780014235  
linux : 0.009722364676153268  
exo : 0.009395714416167291  
monitor : 0.009296862889490732  
snapchat : 0.008726528030189203

[Topic 2]

twitter : 0.09948925415844262  
tag : 0.0735915278124991  
ifttt : 0.07187079477700534  
rt : 0.031373806155106945  
tech : 0.02957588264737773  
trend : 0.02553501737342278  
canada : 0.023853412077789835  
wallpaper : 0.023646846821076786  
central : 0.01639424999864219

sky : 0.01453942024849919  
authority : 0.013805131854674269  
pocketnow : 0.013249517708734315  
tweet : 0.011512189573575934  
androidauth : 0.011374590301579997  
phandroid : 0.010928041369318013  
facebookpages : 0.010259588275808269  
sunset : 0.008963168392543094  
review : 0.008735166248242431  
literally : 0.007919459969613046  
fact : 0.007045931190551231

#### [Topic 3]

sm : 0.047780676897771344  
gb : 0.04319419927614894  
black : 0.04154046246260742  
unlock : 0.03313379170419474  
midnight : 0.030762633924182325  
buy : 0.021442947560195318  
new : 0.0185181812707691  
hot : 0.01828007736677942  
cell : 0.017397962091330457  
smartphone : 0.016939770356931796  
gray : 0.015954727478019994  
plus : 0.015169571043869342  
usa : 0.013381793164040492  
gold : 0.012228758265259288  
hardwarezone : 0.011923305911720295  
silver : 0.011901509917271411  
factory : 0.01106312166319571  
verizon : 0.01065469502581019  
blue : 0.01029271379056016  
edition : 0.00951095237791331

#### [Topic 4]

gb : 0.037486592054155644  
ram : 0.021726597125692092  
protector : 0.017953590075589426  
nokia : 0.01718515443263322  
camera : 0.014914665052254007  
snapdragon : 0.01461055814721028  
dual : 0.013311941763593441  
leak : 0.012460963015819091  
active : 0.012069960365180299  
oneplus : 0.012042996803531601  
storage : 0.011899592616774386  
price : 0.009806549826336946  
glass : 0.009771624394529214  
processor : 0.009660444310400623  
sim : 0.009451270598235676  
screen : 0.00926240709046109  
exynos : 0.008959732435901518  
qualcomm : 0.008744810961050797  
variant : 0.008638023488809774  
core : 0.008193271483404041

#### [Topic 5]



vs : 0.03732839682019339  
tag : 0.03724716307955657  
plus : 0.03677043111377488  
review : 0.027966582731792967  
samsunggalaxy : 0.02790934906672396  
video : 0.025052012817627683  
tech : 0.024406775087481437  
news : 0.02372713099100591  
google : 0.02100027504386046  
unboxing : 0.0206372677433747  
unboxyourphone : 0.020268909665089384  
flag : 0.019463355696413682  
lg : 0.01865277569126849  
irrelevant : 0.01865003478663751  
youtube : 0.018225979640022347  
gadget : 0.017918126872637528  
smartphone : 0.016783183650173766  
comparison : 0.01659543457594826  
crash : 0.016574982580934348  
phone : 0.01580066723909783

[Topic 6]

case : 0.06916029000863662  
charge : 0.020825427868707277  
plus : 0.01983616639631788  
cover : 0.019186817704493882  
work : 0.01864295342132865  
charger : 0.015555918883688544  
fit : 0.012587637551242429  
cable : 0.010544185366986931  
protection : 0.010452880336884114  
wireless : 0.010160425267176277  
phone : 0.009975415122745981  
great : 0.009903877282459634  
drop : 0.009891423120508697  
leather : 0.009634579866624528  
spigen : 0.00952876483793333  
fast : 0.009238725967594189  
protect : 0.008935015180624438  
orchid : 0.008620234042130723  
good : 0.008144529089332716  
tpu : 0.007981544949822865

[Topic 7]

phone : 0.009515392409287065  
new : 0.00859758154970041  
display : 0.005731316122863418  
launch : 0.005715067536826397  
screen : 0.00550505426784757  
device : 0.005393550226497029  
bixby : 0.005326982797831303  
note : 0.005305034546894121  
feature : 0.0049290264482895465  
1 : 0.004754393049320687  
flagship : 0.0046087435141167555  
make : 0.004385665366805055  
smartphone : 0.004299481713947154

like : 0.0042276195988919015  
come : 0.004179053585469187  
look : 0.004160609016774414  
year : 0.00408333136872767  
camera : 0.0039859590708253515  
just : 0.003921727729671118  
good : 0.0038409114607552727

[Topic 8]

love : 0.02230886253618591  
otterbox : 0.01935389154278483  
amaze : 0.018622363568244083  
photography : 0.018495260591904313  
shoot : 0.018322635379162056  
photo : 0.016716859196319964  
instagood : 0.013504672947735526  
spring : 0.012753208340096134  
recovery : 0.01221458020534909  
photooftheday : 0.011328562795161048  
nature : 0.011325147818411752  
tutorial : 0.01129770992154491  
cute : 0.011131808759521886  
beautiful : 0.010392023995140715  
tag : 0.010240277712328112  
samsunggalaxy : 0.009447937630440122  
picture : 0.00929289084685844  
lol : 0.009273849541039868  
happy : 0.00878375647831126  
nofilter : 0.008709405205897575

[Topic 9]

app : 0.014934176382772875  
use : 0.009283640279277566  
button : 0.009096391222315255  
work : 0.007235155303009214  
bixby : 0.006658602120881599  
phone : 0.006391407875892996  
download : 0.0055541712016843985  
user : 0.005217064341355947  
play : 0.00512353173012694  
google : 0.004778909594207152  
update : 0.004741245260947795  
iris : 0.004574852281995002  
new : 0.004500541659447313  
try : 0.004368627323125851  
music : 0.0043291356212816715  
recognition : 0.004311893852065395  
game : 0.004291029451961403  
scanner : 0.004032503766164054  
microsoft : 0.003947638253331034  
message : 0.003914852082678205

## Experimentation with bigrams (instead of unigrams as per above)

```
In [0]: from sklearn.feature_extraction.text import CountVectorizer
countvectorizer = CountVectorizer(min_df=50, stop_words=my_stop_words, m
ax_df=0.8, tokenizer=word_tokenizer, ngram_range=(2,2), max_features=100
00)
count_matrix = countvectorizer.fit_transform(corpus)
print('Number of features in tf-idf:', len(countvectorizer.vocabulary_))
# a mapping of terms to feature indices
```

Number of features in tf-idf: 10000

```
In [0]: count_matrix = count_matrix.toarray()
sum_bigrams = count_matrix.sum(axis=0)
words_freq = [(word, sum_bigrams[idx]) for word, idx in countvectorizer.
vocabulary_.items()]
words_freq = sorted(words_freq, key = lambda x: x[1], reverse=True)
```

```
In [0]: words_freq[:30]
```

```
Out[0]: [('pre order', 13278),
('gb ram', 8942),
('t mobile', 8295),
('appear \uffffl', 7467),
('infinity display', 6357),
('new phone', 5603),
('new flagship', 5374),
('gear vr', 5336),
('screen protector', 4944),
('google pixel', 4895),
('home button', 4839),
('tag ifttt', 4769),
('iris scanner', 4667),
('pixel xl', 4478),
('bixby button', 4340),
('look like', 4275),
('fingerprint sensor', 4090),
('south korea', 3956),
('bixby voice', 3814),
('ram gb', 3799),
('flag irrelevant', 3799),
('new york', 3777),
('flagship phone', 3642),
('facial recognition', 3582),
('midnight black', 3522),
('aspect ratio', 3505),
('battery life', 3405),
('wireless charge', 3336),
('red tint', 3260),
('fingerprint scanner', 3254)]
```

## Most important attributes for Galaxy S8:

The 30 most frequently occurring bigrams in the reviews/post based on the above analysis. This gives us a hint on what features/characteristics of the phone are being discussed by the users the most. Some of the most talked about attributes based on our analysis above were:

1. Memory (as indicated by 'gb ram')
2. Infinity display
3. Iris Scanner
4. Bixby button
5. Fingerprint sensor/scanner
6. Bixby voice
7. Facial recognition
8. Red tint (an issue faced by some users with Samsung products)
9. Battery life
10. Wireless charger

In [0]: